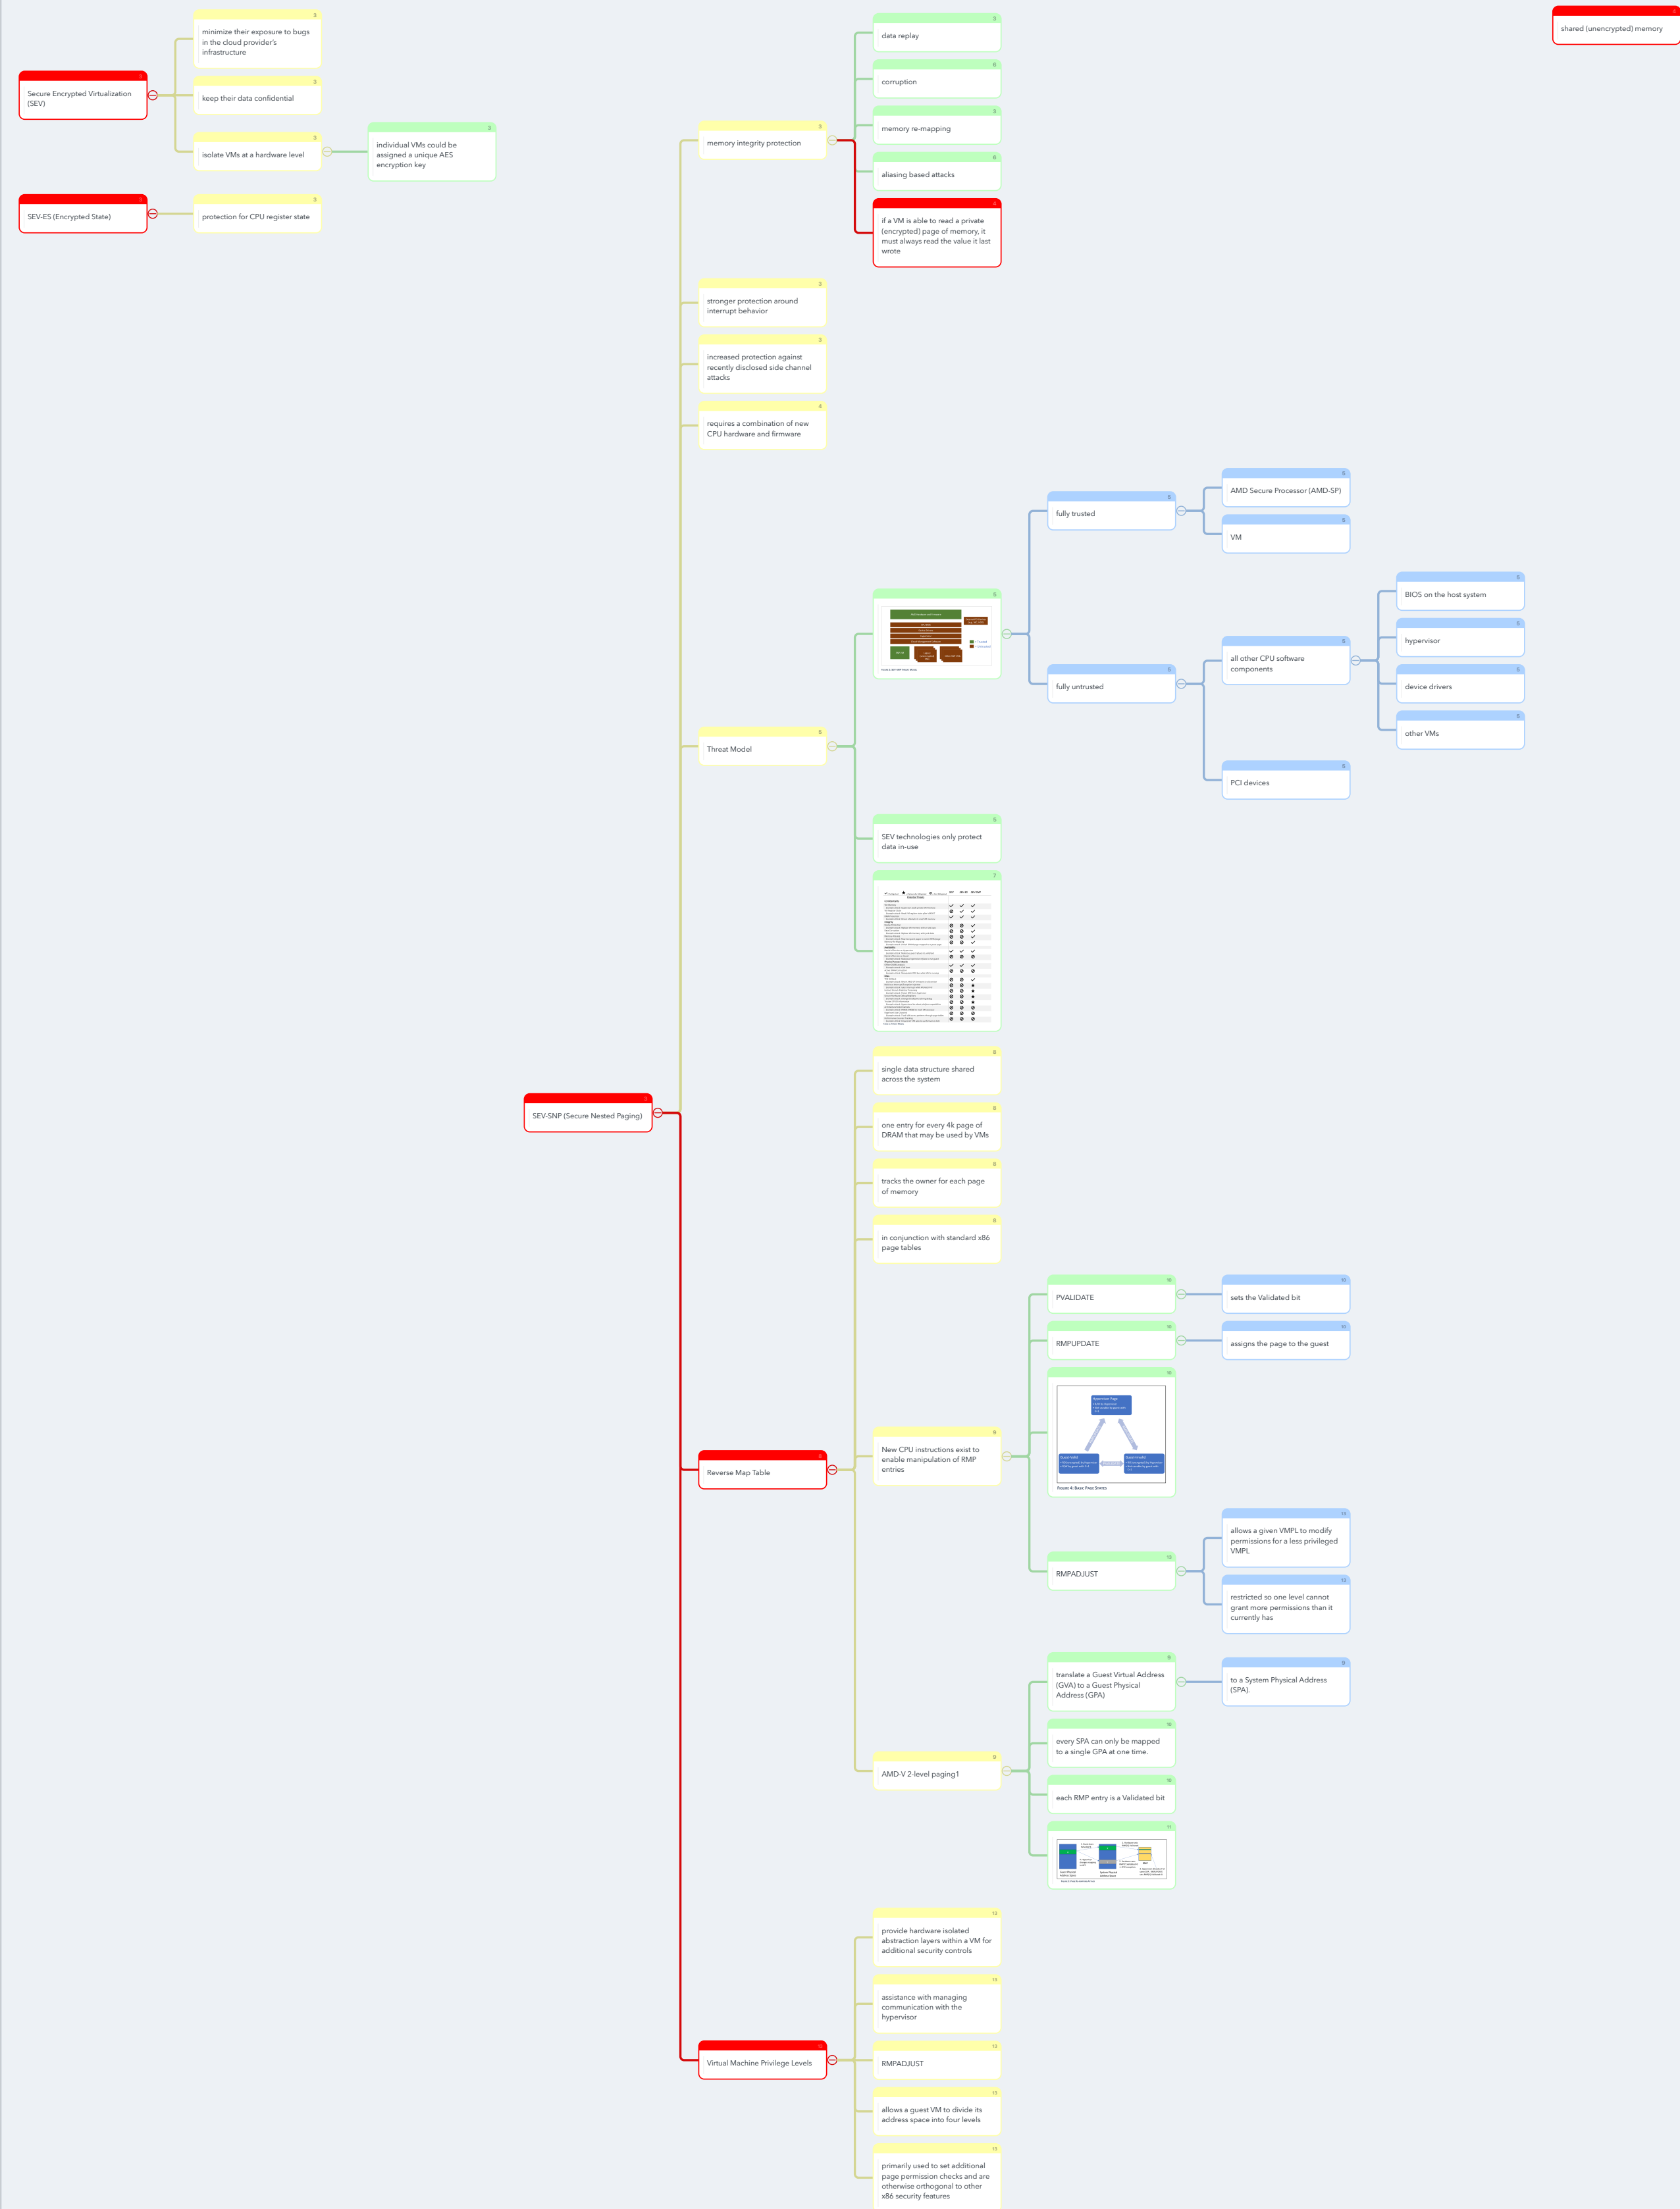


SEV-SNP-
strengthening-vm-
isolation-with-
integrity-protection-
and-
more-1e2d0f1f63c91
68a575bd487c8c869
8b



Secure Encrypted Virtualization (SEV)	3
minimize their exposure to bugs in the cloud provider’s infrastructure	3
keep their data confidential	3
isolate VMs at a hardware level	3
individual VMs could be assigned a unique AES encryption key	3
SEV-ES (Encrypted State)	3
protection for CPU register state	3
SEV-SNP (Secure Nested Paging)	3
memory integrity protection	3
data replay	3
corruption	6
memory re-mapping	3
aliasing based attacks	6
if a VM is able to read a private (encrypted) page of memory, it must always read the value it last wrote	4
stronger protection around interrupt behavior	3
increased protection against recently disclosed side channel attacks	3
requires a combination of new CPU hardware and firmware	4
Threat Model	5

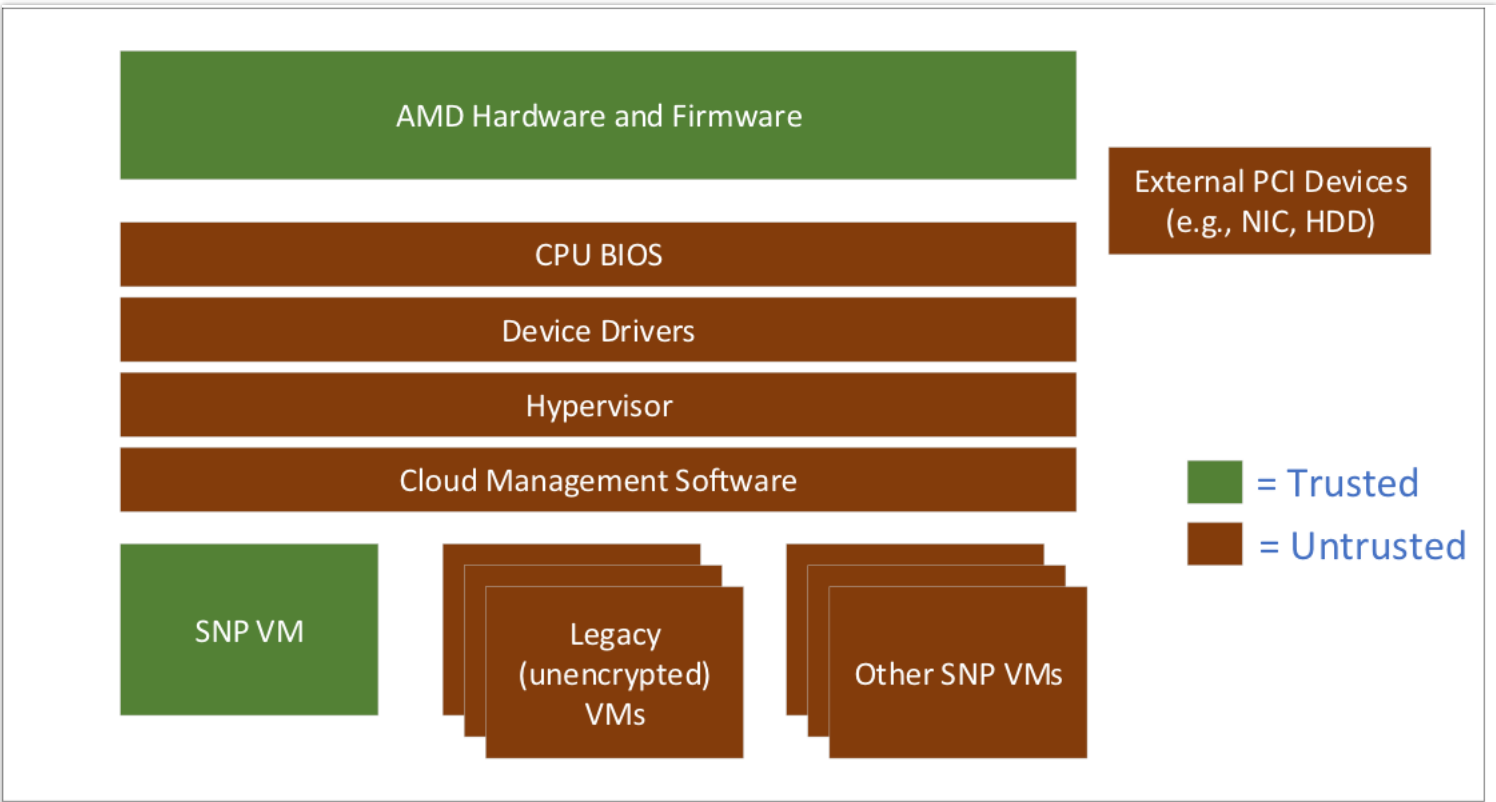


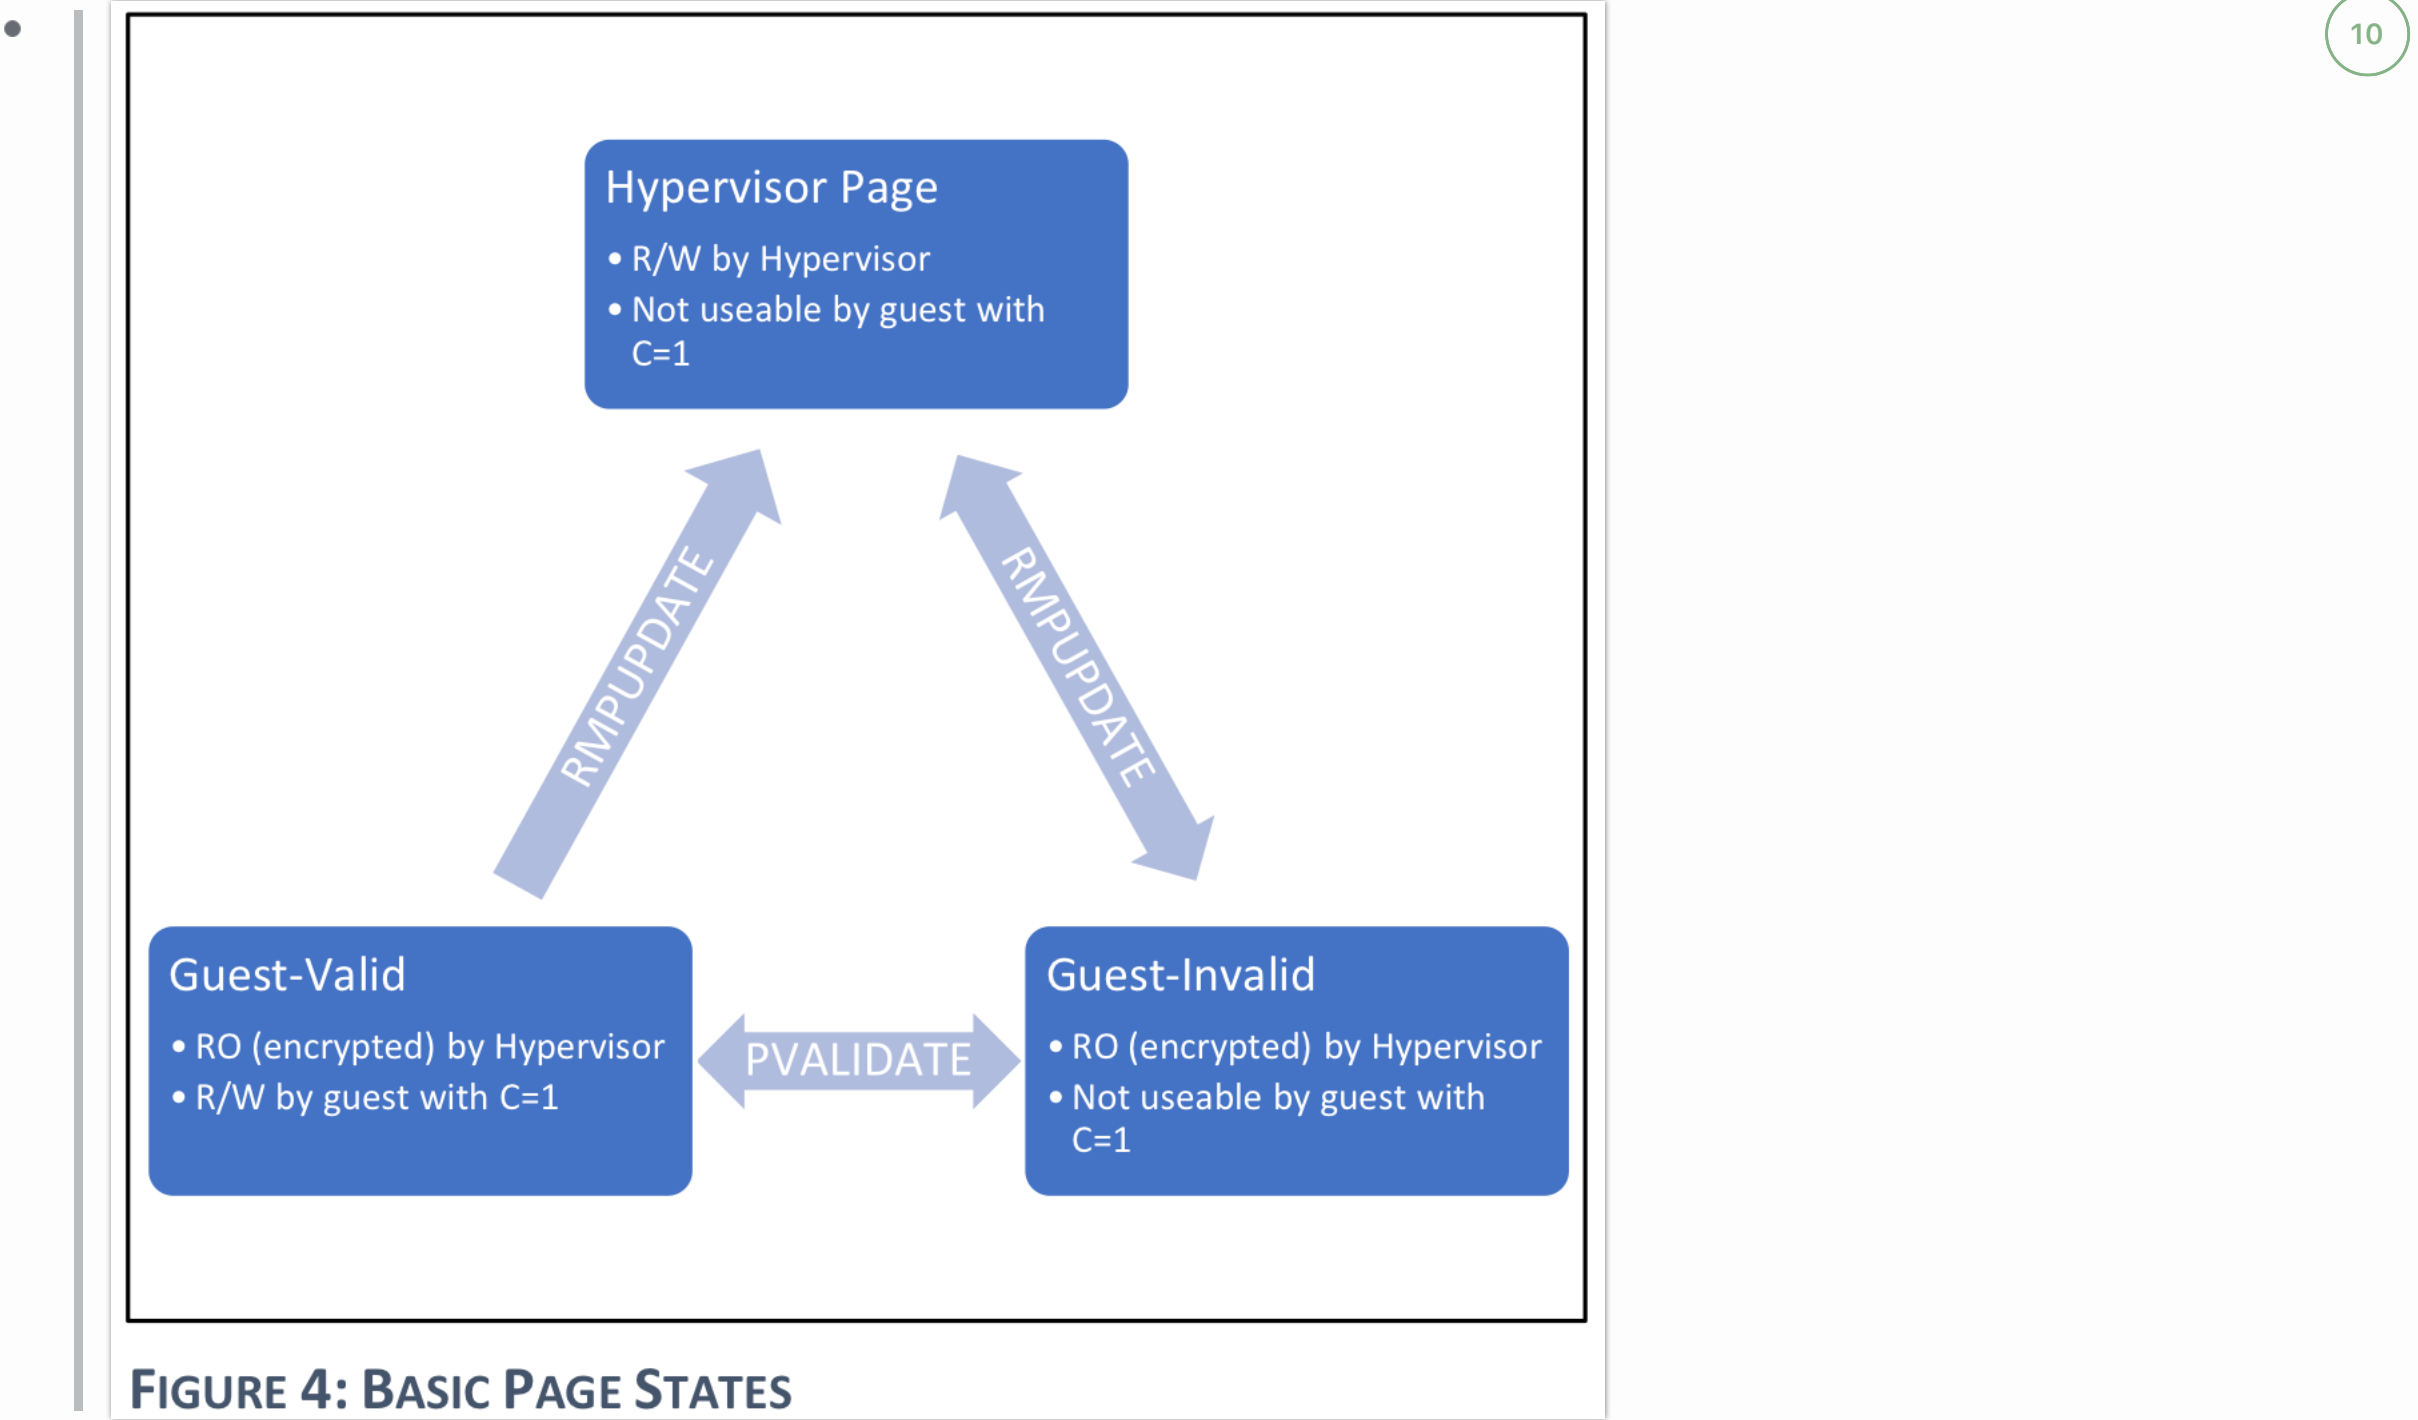
FIGURE 2: SEV-SNP THREAT MODEL

<div> <div>▼</div> <div>fully trusted</div> </div>	5
<div> <div>•</div> <div>AMD Secure Processor (AMD-SP)</div> </div>	5
<div> <div>•</div> <div>VM</div> </div>	5
<div> <div>▼</div> <div>fully untrusted</div> </div>	5
<div> <div>▼</div> <div>all other CPU software components</div> </div>	5
<div> <div>•</div> <div>BIOS on the host system</div> </div>	5
<div> <div>•</div> <div>hypervisor</div> </div>	5
<div> <div>•</div> <div>device drivers</div> </div>	5
<div> <div>•</div> <div>other VMs</div> </div>	5
<div> <div>•</div> <div>PCI devices</div> </div>	5
<div> <div>•</div> <div>SEV technologies only protect data in-use</div> </div>	5

<div> <div>•</div> <div> <div> <div>✓ = Mitigated</div> <div>★ = Optionally Mitigated</div> <div>⊘ = Not Mitigated</div> </div> <div> <div>SEV</div> <div>SEV-ES</div> <div>SEV-SNP</div> </div> <div> <div>Potential Threats</div> <div> <div>Confidentiality</div> <div> <div>VM Memory</div> <div>Example attack: Hypervisor reads private VM memory</div> </div> <div> <div>VM Register State</div> <div>Example attack: Read VM register state after VMEXIT</div> </div> <div> <div>DMA Protection</div> <div>Example attack: Device attempts to read VM memory</div> </div> <div>Integrity</div> <div> <div>Replay Protection</div> <div>Example attack: Replace VM memory with an old copy</div> </div> <div> <div>Data Corruption</div> <div>Example attack: Replace VM memory with junk data</div> </div> <div> <div>Memory Aliasing</div> <div>Example attack: Map two guest pages to same DRAM page</div> </div> <div> <div>Memory Re-Mapping</div> <div>Example attack: Switch DRAM page mapped to a guest page</div> </div> <div>Availability</div> <div> <div>Denial of Service on Hypervisor</div> <div>Example attack: Malicious guest refuses to yield/exit</div> </div> <div> <div>Denial of Service on Guest</div> <div>Example attack: Malicious hypervisor refuses to run guest</div> </div> <div>Physical Access Attacks</div> <div> <div>Offline DRAM analysis</div> <div>Example attack: Cold boot</div> </div> <div> <div>Active DRAM corruption</div> <div>Example attack: Manipulate DDR bus while VM is running</div> </div> <div>Misc.</div> <div> <div>TCB Rollback</div> <div>Example attack: Revert AMD-SP firmware to old version</div> </div> <div> <div>Malicious Interrupt/Exception Injection</div> <div>Example attack: Inject interrupt while RFLAGS.IF=0</div> </div> <div> <div>Indirect Branch Predictor Poisoning</div> <div>Example attack: Poison BTB from hypervisor</div> </div> <div> <div>Secure Hardware Debug Registers</div> <div>Example attack: Change breakpoints during debug</div> </div> <div> <div>Trusted CPUID Information</div> <div>Example attack: Hypervisors lies about platform capabilities</div> </div> <div> <div>Architectural Side Channels</div> <div>Example attack: PRIME+PROBE to track VM accesses</div> </div> <div> <div>Page-level Side Channels</div> <div>Example attack: Track VM access patterns through page tables</div> </div> <div> <div>Performance Counter Tracking</div> <div>Example attack: Fingerprint VM apps by performance data</div> </div> </div> </div> </div> </div>	7
---	---

TABLE 1: THREAT MODEL

▼ Reverse Map Table	8
• single data structure shared across the system	8
• one entry for every 4k page of DRAM that may be used by VMs	8
• tracks the owner for each page of memory	8
• in conjunction with standard x86 page tables	8
▼ New CPU instructions exist to enable manipulation of RMP entries	9
▼ PVALIDATE	10
• sets the Validated bit	10
▼ RMPUPDATE	10
• assigns the page to the guest	10



▼ RMPADJUST	13
• allows a given VMPL to modify permissions for a less privileged VMPL	13
• restricted so one level cannot grant more permissions than it currently has	13
▼ AMD-V 2-level paging1	9

▼ | translate a Guest Virtual Address (GVA) to a Guest Physical Address (GPA)

9

• | to a System Physical Address (SPA).

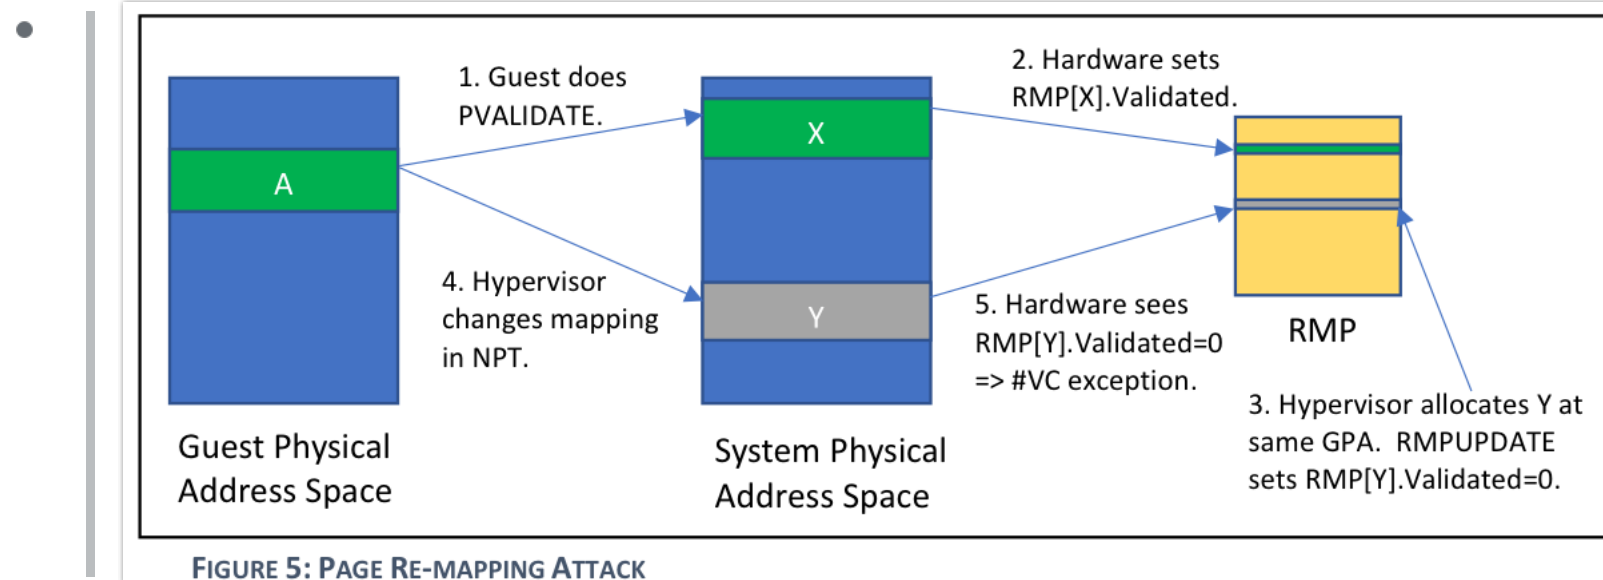
9

• | every SPA can only be mapped to a single GPA at one time.

10

• | each RMP entry is a Validated bit

10



11

▼ | Virtual Machine Privilege Levels

13

• | provide hardware isolated abstraction layers within a VM for additional security controls

13

• | assistance with managing communication with the hypervisor

13

• | RMPADJUST

13

• | allows a guest VM to divide its address space into four levels

13

• | primarily used to set additional page permission checks and are otherwise orthogonal to other x86 security features

13

• | shared (unencrypted) memory

4